

COMP6033 - Individual Research Report

Henry Lovett

Abstract—The abstract goes here.

Index Terms—IEEEtran, journal, L^AT_EX, paper, template.

I. INTRODUCTION

The desire for low power devices has been driven by the mobile age. Companies are competing on battery life of portable devices, such as smartphones and tablets. This drive for low power has resulted in different synthesis techniques. This paper will review and explore some of the techniques used to help reduce the power consumption of a circuit, with particular interest on the synthesis techniques used to do so.

There are two main categories of power consumption - dynamic and leakage. Dynamic power is the power consumed when the circuit is enabled and functioning. Power is used here through the charging and discharging of the internal capacitors of the module. Leakage power is due to the non-ideal characteristics of sub-micron CMOS transistors. Leakage power mainly occurs when the module is in a idle, non-switching state [1] and can contribute around 22% of the total power consumption of a 90nm FPGA [2]

The report begins with a review of different techniques in turn. This includes a brief introduction to the theory and a discussion of synthesis problems and solutions. The report concludes by reviewing all techniques and their relevant advantages and disadvantages.

II. POWER TECHNIQUES

A. Power Gating

Leakage currents begin to occur in sub-micron technologies, typically lower than 100nm [1], [3]. With each technology generation, gate leakage increases 30 fold [4]. Leakage currents can be identified as two sources - subthreshold currents and gate tunnelling. Subthreshold currents occur when the voltage applied to the gate is lower than the transistor threshold voltage. The channel of the transistor during weak inversion still allows some current to flow from source to drain. This current is the subthreshold current and occurs when the transistor is off and is also exponentially dependant on V_{th} [5]. Gate tunnelling currents increase with thinner oxide layer [6]. A thin oxide results in electrons being able to pass through the gate into the substrate of the transistor. This thin gate results in electrons passing through it, into the circuit causing leakage. Gate tunnelling occurs when the transistor is active [3].

1) *Theory*: Power gating, in principle, is where the power to a module is switched off when not in use. By doing this, the module does not consume any power. It can be achieved by using either a header or a footer switch to disable either the supply connection or the ground connection. Figure 1 shows the realisation of the power gating circuits.

Although the theory of operation is simple, the technique poses many issues in the implementation. Firstly, as the module is floating, the outputs are undefined. This




Fig. 1. Circuit diagrams showing the use of footer switches (left) or header switches (right).

Break this into two subfigures

can cause the gates in another module to short circuit - at an input voltage of $V_{dd}/2$, both the NMOS and PMOS transistors will be on, resulting in a direct line from supply to ground. The solution to this is to add logic gates with a 'clamp' signal. This could be either an AND gate, of which the clamp signal is active low, or an OR gate with an active high clamp. This added logic is needed per signal output of the power gated module.

The second issue that power gating raises is when the gated module contains sequential logic. By removing the power to the sequential logic, the state is lost. This can then cause issues as state retention may be necessary, as well as low power. The solution is to use a state retention register. There is a timing overhead involved to store the register before putting the device to sleep, disabling the majority of the circuit. State retention registers also require an individual power supply, meaning all power gated modules require two supplies; one which is gated and one constant supply.

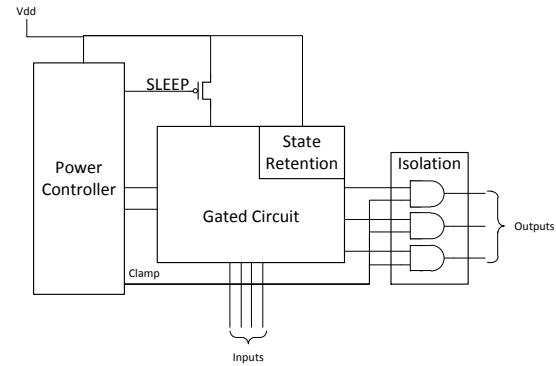


Fig. 2. The realisation of a power gated module. State retention, isolation and a power manager are all needed

The final aspect of power gating that needs addressing is the power management. The power manager is responsible for when to turn the module on and off. As well as the control of the sleep transistors, the power manager is also responsible for the stopping the clock, clamping the outputs and saving the state to the state retention mechanism.

There is a large overhead in the additional circuitry needed for power gating. The main increase are the state retention registers as these can require 10% more area in silicon [?]

Citation needed

When footer switches are used, the voltage that at the source of the NMOS is known as the virtual ground (VGND). When the footer is active, $VGND \approx GND$. When off, $VGND \approx V_{dd}$. A full realisation of a power gated circuit is shown in figure 2.

2) *Synthesis Techniques*: Fine and coarse grain techniques Power gating can be split into two techniques, fine and coarse grain. Fine grain gating is where individual cells (e.g. in ASIC) have their own gating transistor. Coarse grain techniques are where gating transistors are applied on a larger module scale. [3] compares the

fine grain and coarse grain methods on a 65nm FPGA technology. Here, a SRAM block is used as the test circuit. Fine grain, in this case, is where the memory of each cell is individually gated, and coarse is where a 16×1 SRAM module is gated. It was conclusively shown that coarse grain gating reduced power consumption more than fine grain due to being able to turn off the read/write circuitry.

When deciding to sleep a module, there are a number of overheads. Due to the charging / discharging of the gated circuit (depending on the use of header or footer switches), when the circuit is re-enabled, there is an amount of time needed to allow the capacitors to return to normal [7]. This is referred to as the wake-up time and can also include the time to restore state if applicable. The (dis)charging can also poses ground bounce issues to the device due to the sudden draw of current, [8], [9].

This wake up time can then impact the performance of the device - if more energy is consumed waiting for the disabled module than it would have consumed being awake, the overall power is not reduced. [10] discusses a method where an intermediate power saving mode is introduced. This is done by using a standard footer power gating NMOS. An extra PMOS is then added in parallel with the power gating switch. Table I summarises the states of the transistors in the circuit in figure 3

The RUN and COLD states operate as previously discussed. The PARK state, however, is where the virtual ground voltage does not rise to V_{dd} due to the PMOS conducting. The virtual ground is at the threshold voltage of the PMOS, so the leakage current is less than in the RUN state, but not so high that state retention is needed. This drastically reduces the wake up overhead of the circuit.

A further improvement to this single extra state is proposed in [11]. Here, multiple intermediate power states




Fig. 3. Middle power state implementation used in [10]

TABLE I
POWER SAVING MODES OF [10]

State Name	NMOS State	PMOS State	Virtual Ground Voltage (V)
RUN	ON	OFF	GND
COLD	OFF	OFF	V_{DD}
PARK	OFF	ON	V_{th} PMOS

are implemented by using different subthreshold gate voltages on the NMOS footer switch. The compromise of wake up time against leakage reduction can then be more finely controlled to best suit the circuit. In general, the closer the virtual ground is to ground, the more leakage current occurs, but a faster wake up is available.

B. Power Scaling

I want to try and find use of two modules, one with a Low V_t for high performance, high power, and a high V_t for low performance, low power, which can be switched between. If not, just look at multi V_t technologies

III. FREQUENCY DOMAINS

Lorem Ipsum...

IV. FREQUENCY SCALING

Lorem Ipsum...




Fig. 4. Clock gating circuit using an AND gate

Gated Module or Gated Circuit? Consistency with the other diagrams




Fig. 5. Clock gating circuit using an AND gate and a latch

Gated Module or Gated Circuit? Consistency with the other diagrams

V. CLOCK GATING

A. Theory

The clock in a sequential circuit can contribute 15-45% of the power [12]. Therefore it is a large area of potential power saving. Clock gating is an approach of controlling the clock to individual modules of a design by either stopping or slowing down the clock with respect to a master clock [13]. An approach, seen in [14], involves stopping the clock to unused modules.

This method can be realised using two simple circuits seen in figures 4 and 5. Although figure 4 is functionally complete, in reality, a latch is needed to remove any glitches in the circuit. These are fundamentally different to load-enable registers, where the input is multiplexed between the current value or the input. The load-enable registers are still clocked at the master clock frequency.

B. Synthesis Techniques

A gating function is typically defined by the designer within the RTL design stage. However, a more common approach is to allow the synthesis tool to obtain the gating functions from a gate-level netlist [15], [16].

The general outline for the synthesis is to find the

clock gating function for each flip-flop. The flip-flops are then grouped so that they are driven by the same function. The problem of simplifying the gating function is looked at in [17]. Here, an algorithm is suggested where the gating function is shared by existing combinational logic. This was shown to reduce the logic added by introducing clock gating.

However, sometimes the addition of clock gating is not advantageous. The gating function can be large and therefore can cause timing violations, resulting in an unsuitable synthesis. If the gating function is large enough, it can also consume more power than it saves. Both of these issues are addressed in [18]. Here, the author proposed solutions to large gating functions by reducing the depth of logic by an approximation. The approximation is made such that the resulting logic does not gate the clock more than the original function. It results in the flip-flop being clocked more often, but reduces the logic so that it can be utilised, thereby saving some power.

[18] also proposes the use of a clustering algorithm. The algorithm looks at grouping similar gating functions. This can then reduce the overall logic needed to implement the clock gating and maximise the energy saved.

C. Conclusion

Clock gating is a simple principle to implement on a small scale. The underlying theory is to disable a module when it is not in use. This is done by identifying a gating function which disallows clock propagation if the function is asserted.

However, clock gating can produce large gating functions which violate the timing constraints of the circuit, or even consume more power than they save. This results in two problems that need solving - group formation and simplification [17], [19].

mds

17th March, 2014

REFERENCES

- [1] A. A. Bsoul and S. J. Wilton, "An fpga architecture supporting dynamically controlled power gating," in *Field-Programmable Technology (FPT), 2010 International Conference on*. IEEE, 2010, pp. 1–8.
- [2] Altera, "Statix ii vs. virtex-4 power comparison & estimation accuracy," 2005, online at http://www.altera.co.uk/literature/wp/wp_s2v4_pwr_acc.pdf.
- [3] P. S. Nair, S. Koppa, and E. B. John, "A comparative analysis of coarse-grain and fine-grain power gating for fpga lookup tables," in *Circuits and Systems, 2009. MWSCAS'09. 52nd IEEE International Midwest Symposium on*. IEEE, 2009, pp. 507–510.
- [4] K. Bernstein, C.-T. Chuang, R. Joshi, and R. Puri, "Design and cad challenges in sub-90nm cmos technologies," in *Computer Aided Design, 2003. ICCAD-2003. International Conference on*. IEEE, 2003, pp. 129–136.
- [5] S. Borkar, "Design challenges of technology scaling," *Micro, IEEE*, vol. 19, no. 4, pp. 23–29, 1999.
- [6] W. M Arden, "The international technology roadmap for semiconductors perspectives and challenges for the next 15 years," *Current Opinion in Solid State and Materials Science*, vol. 6, no. 5, pp. 371–377, 2002.
- [7] A. Abdollahi, F. Fallah, and M. Pedram, "An effective power mode transition technique in mtcmos circuits," in *Proceedings of the 42nd annual Design Automation Conference*. ACM, 2005, pp. 37–42.
- [8] S. Kim, S. V. Kosonocky, and D. R. Knebel, "Understanding and minimizing ground bounce during mode transition of power gating structures," in *Proceedings of the 2003 international symposium on Low power electronics and design*. ACM, 2003, pp. 22–25.
- [9] Y.-S. Chang, S. K. Gupta, and M. A. Breuer, "Analysis of ground bounce in deep sub-micron circuits," in *VLSI Test Symposium, 1997., 15th IEEE*. IEEE, 1997, pp. 110–116.
- [10] S. Kim, S. V. Kosonocky, D. R. Knebel, and K. Stawiasz, "Experimental measurement of a novel power gating structure with intermediate power saving mode," in *Low Power Electronics and Design, 2004. ISLPED'04. Proceedings of the 2004 International Symposium on*. IEEE, 2004, pp. 20–25.
- [11] H. Singh, K. Agarwal, D. Sylvester, and K. J. Nowka, "Enhanced leakage reduction techniques using intermediate strength power gating," *Very Large Scale Integration (VLSI) Systems, IEEE Transactions on*, vol. 15, no. 11, pp. 1215–1224, 2007.
- [12] M. Pedram, "Power minimization in ic design: principles and applications," *ACM Transactions on Design Automation of Electronic Systems (TODAES)*, vol. 1, no. 1, pp. 3–56, 1996.
- [13] Q. Wu, M. Pedram, and X. Wu, "Clock-gating and its application to low power design of sequential circuits," *Circuits and Systems I: Fundamental Theory and Applications, IEEE Transactions on*, vol. 47, no. 3, pp. 415–420, Mar 2000.
- [14] G. E. Téllez, A. Farrahi, and M. Sarrafzadeh, "Activity-driven clock design for low power circuits," in *Proceedings of the 1995 IEEE/ACM international conference on Computer-aided design*. IEEE Computer Society, 1995, pp. 62–65.
- [15] L. Benini, G. De Micheli, E. Macii, M. Poncino, and R. Scarsi, "Symbolic synthesis of clock-gating logic for power optimization of synchronous controllers," *ACM Transactions on Design Automation of Electronic Systems (TODAES)*, vol. 4, no. 4, pp. 351–375, 1999.
- [16] A. P. Hurst, "Automatic synthesis of clock gating logic with controlled netlist perturbation," in *Proceedings of the 45th annual Design Automation Conference*. ACM, 2008, pp. 654–657.
- [17] I. Han and Y. Shin, "Synthesis of clock gating logic through factored form matching," in *IC Design Technology (ICICDT), 2012 IEEE International Conference on*, May 2012, pp. 1–4.
- [18] E. Arbel, C. Eisner, and O. Rokhlenko, "Resurrecting infeasible clock-gating functions," in *Design Automation Conference, 2009. DAC '09. 46th ACM/IEEE*, July 2009, pp. 160–165.
- [19] S. Paik, I. Han, S. Kim, and Y. Shin, "Clock gating synthesis of pulsed-latch circuits," *Computer-Aided Design of Integrated Circuits and Systems, IEEE Transactions on*, vol. 31, no. 7, pp. 1019–1030, July 2012.

Henry Lovett Henry is a fourth year MEng Student at the University of Southampton.